



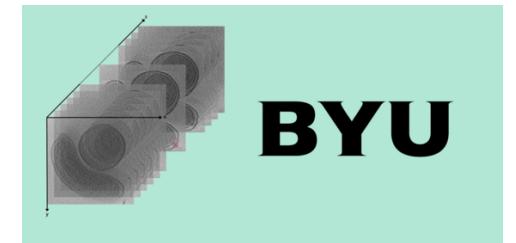
Internship Project Report

Matthew Ward

August 4, 2025

Introduction

- BYU Applied and Computational Math (DS & ML, April 2026)
- Biophysics Simulation Group—computer vision and competition dataset curation with 3D pictures of bacteria (cryo-ET tomograms)
- Data engineer intern with you until August 16th!



Fast, Efficient
3D Imaging

Overview of projects

- 1. Generate a processing performance report**
- 2. Improve registration pipeline with pose graph optimization**
- 3. Model laser spot illumination as a Gaussian from sensor data**

Processing Performance Report

Processing Performance Report

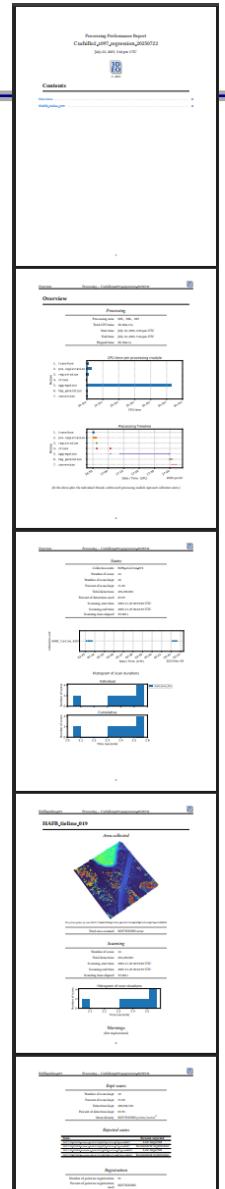
Not to be confused with sensor performance report!

Processing Performance Report

Not to be confused with sensor performance report!

A PDF report summarizing key insights from processing.

Located in <acadia-output-directory>/qc/processing_report after processing is complete



Processing Report Generation Process

Processing directory

for example, albert:/shares/processed/cuchillo/flightData/FlatCreek



JSON of filepaths and statistics

to be used in report



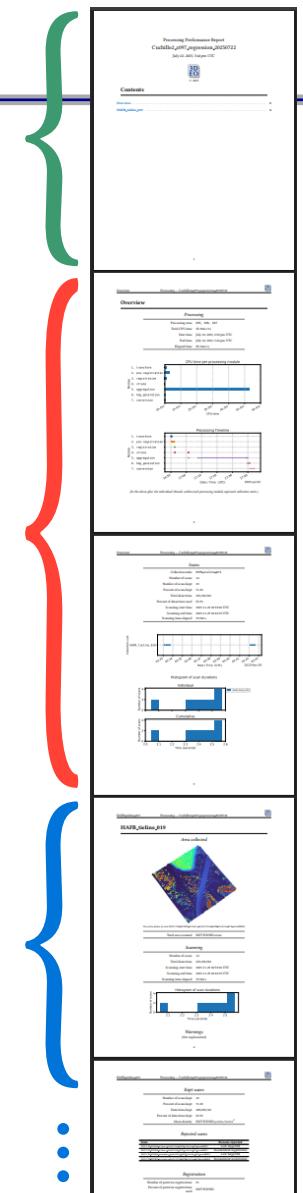
LATEX report



PDF report

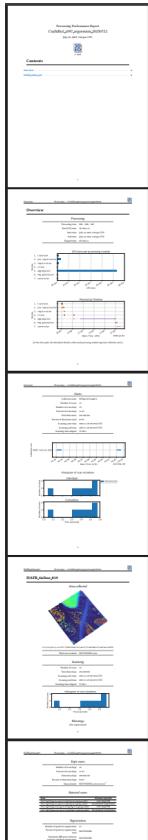
Report Format

- **Contents**
- **Overview** (information about all tiles together)
- **Tile 1** (information specific to this tile)
- **Tile 2**
- :
- **Tile n**



Example Processing Reports

Single Tile Report



(6 pages)

Mapping Report (56 tiles)



(203 pages)

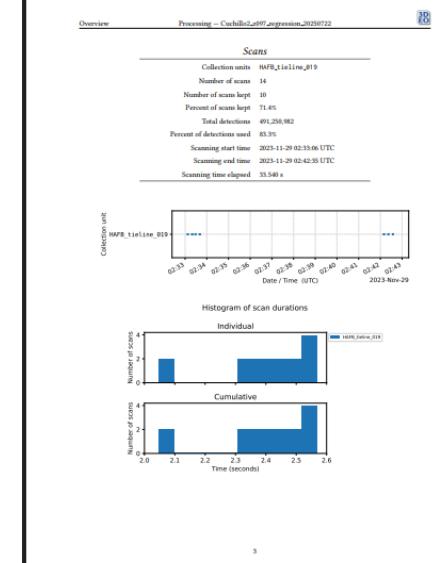
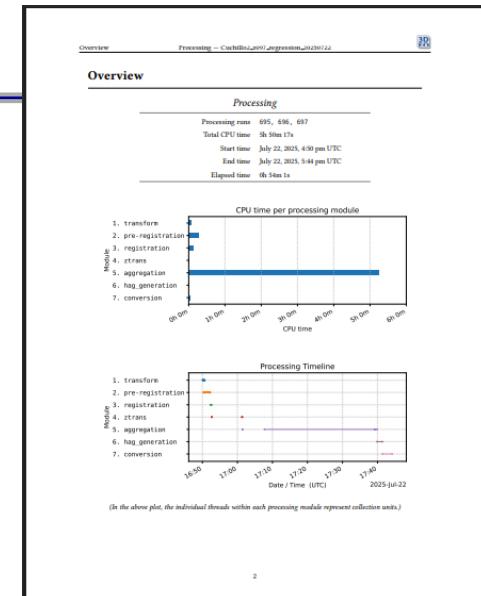
Overview

- Processing summary**

How long processing took, CPU time, etc.

- Scan summary**

Number of detections, how long scanning took, etc.



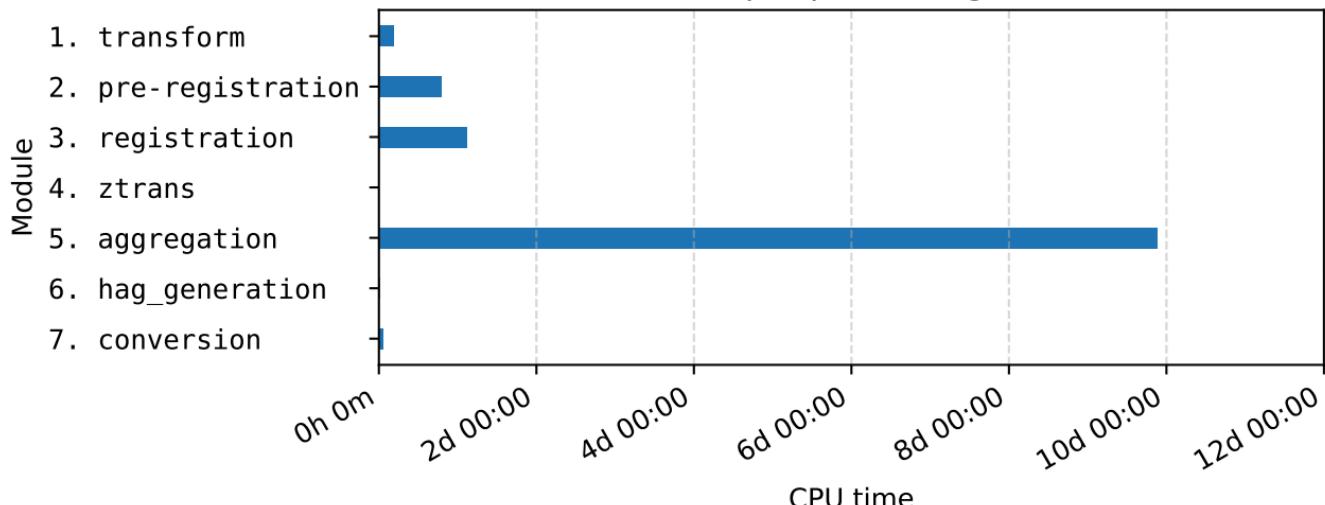
(Overview from Single Tile Report)

Processing Summary

Processing

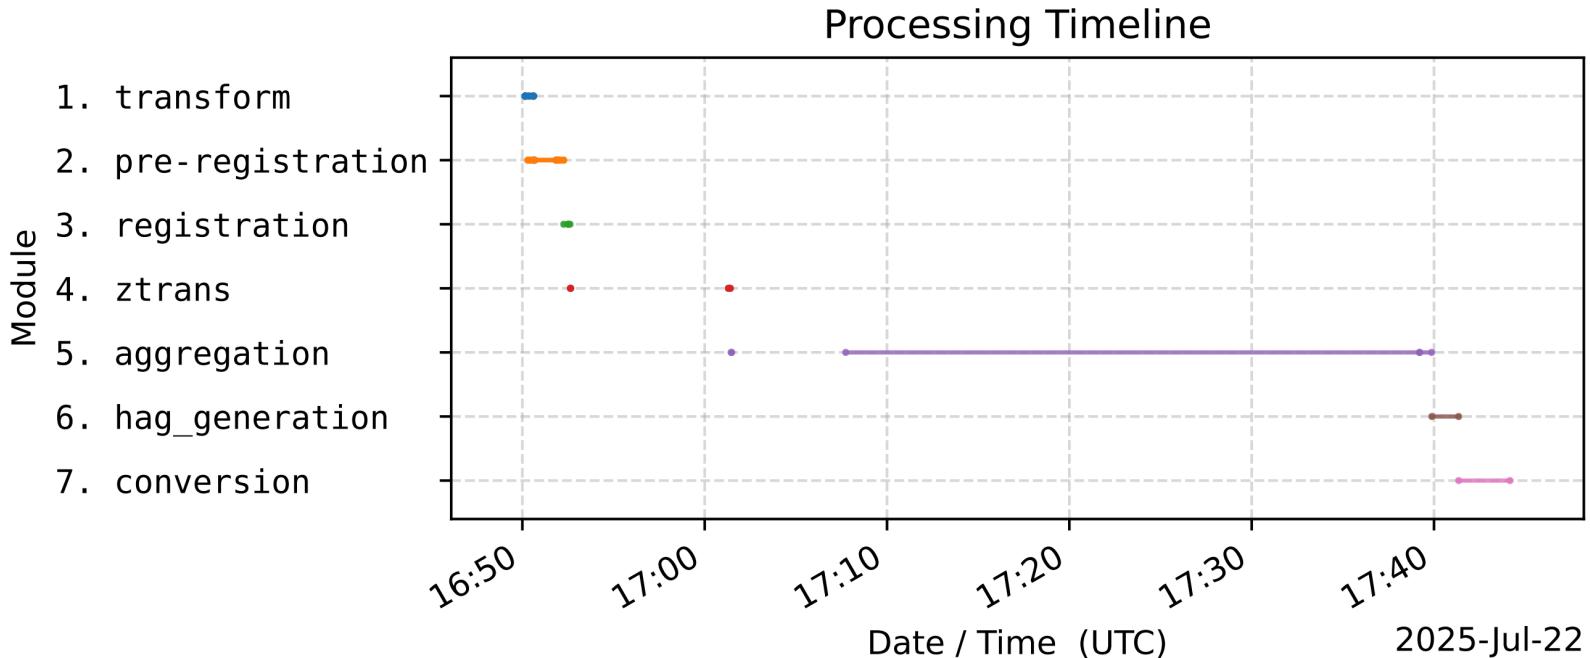
Processing runs	631
Total CPU time	12d 01:30:17
Start time	July 16, 2025, 7:42 pm UTC
End time	July 17, 2025, 5:32 am UTC
Elapsed time	9h 50m 11s

CPU time per processing module

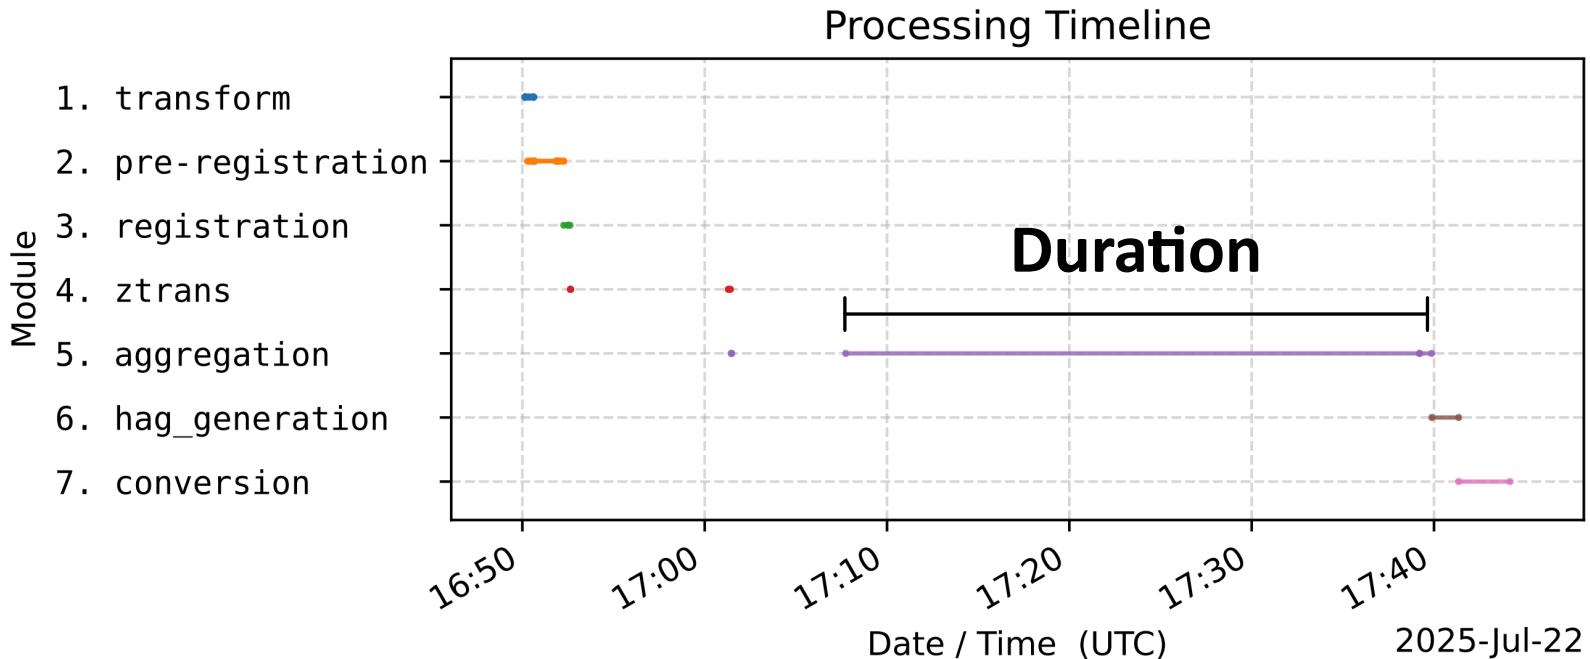


From Mapping Report (56 tiles)

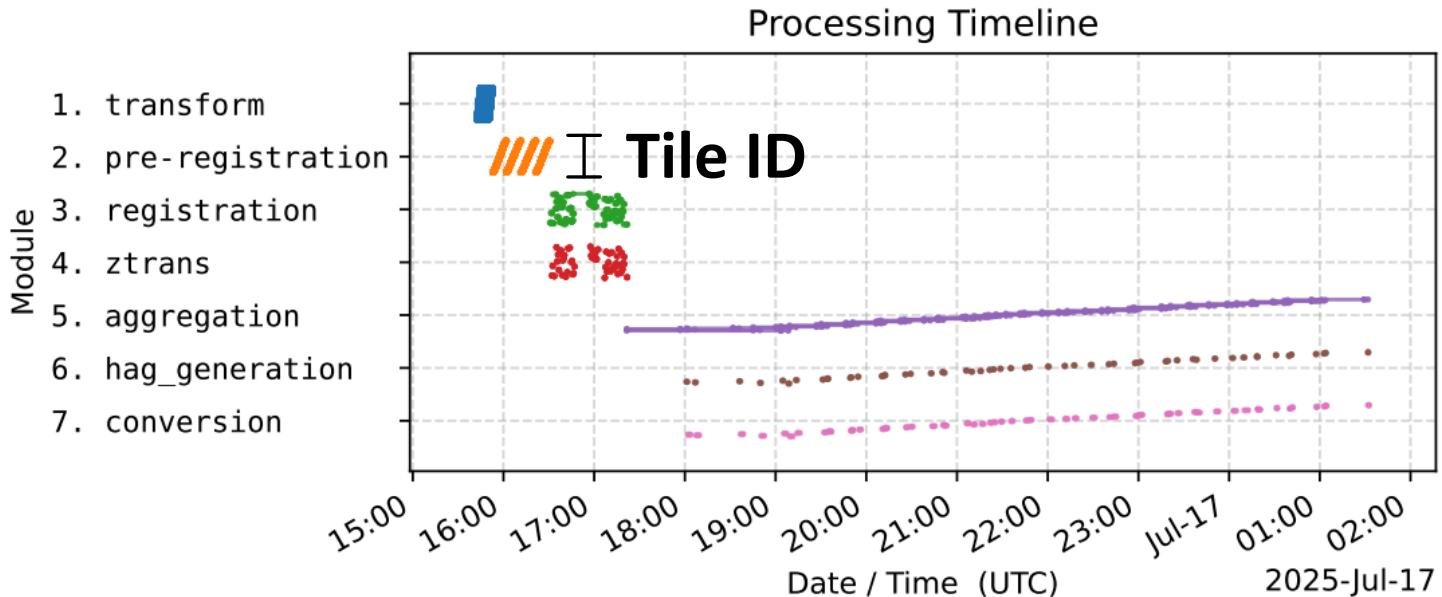
Processing Timeline



Processing Timeline



Processing Timeline



(In the above plot, the individual threads within each processing module represent collection units.)

Per Tile Sections

- **Picture of the tile**
- **Scan information**
 - Number of detections, how long scanning took, etc.
- **Rejected data**
 - Reasons for rejection
- **Registration information**
 - To evaluate success of registration



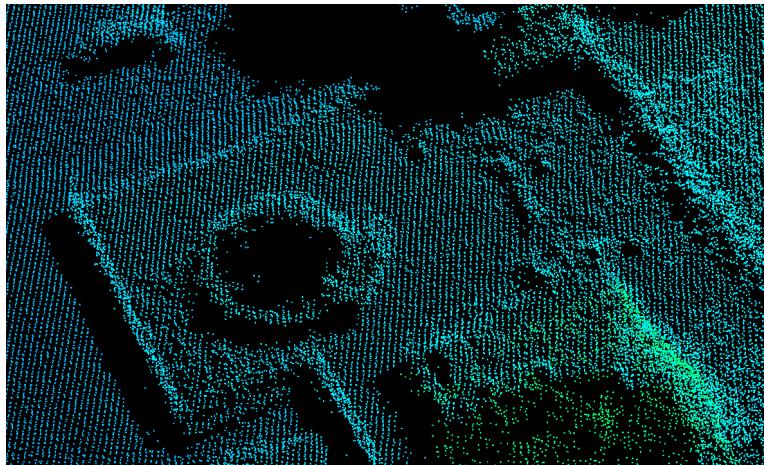
(HAFB_tieline_019 from Single Tile Report)

Future Work

- Continue to refine wording and format for clarity
- Finish implementing fields
 - Scan density information
 - Report on more data rejection reasons
 - Registration success metrics
 - Warnings—for example, warn if the beam was dumping
- Get client feedback—what do they want to see in the report?
- Make \LaTeX compilation container smaller

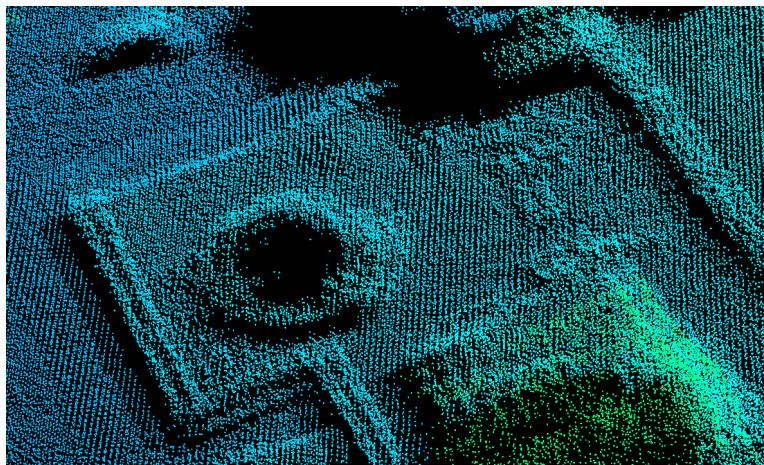
Pose Graphs for Registration

Good registration



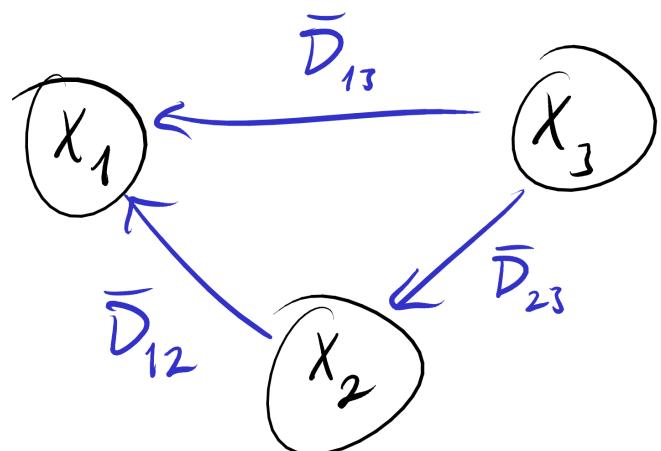
Bad registration

notice that the building is doubled here



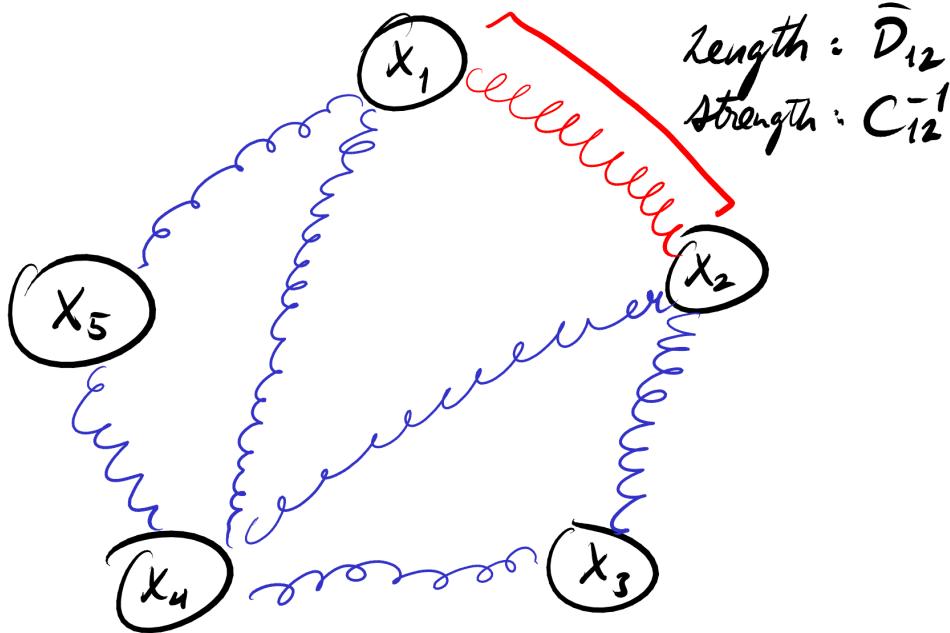
Pose Graph

We have ways to move one scan to align with another, with some uncertainty. How can we move n scans to align with each other?

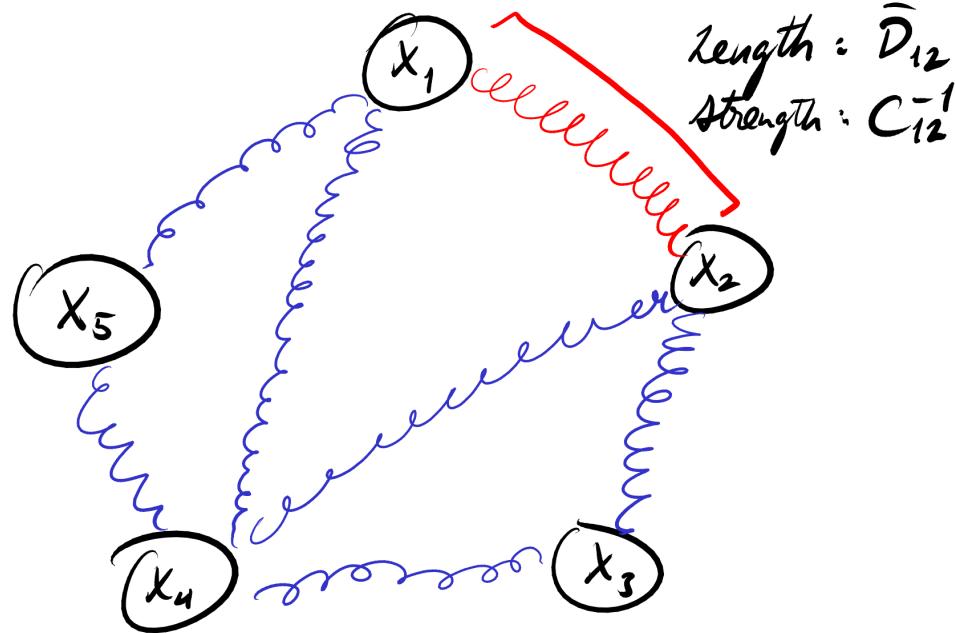


In practice, $\bar{D}_{13} \neq \bar{D}_{12}\bar{D}_{23}$.

Pose Graph as Springs

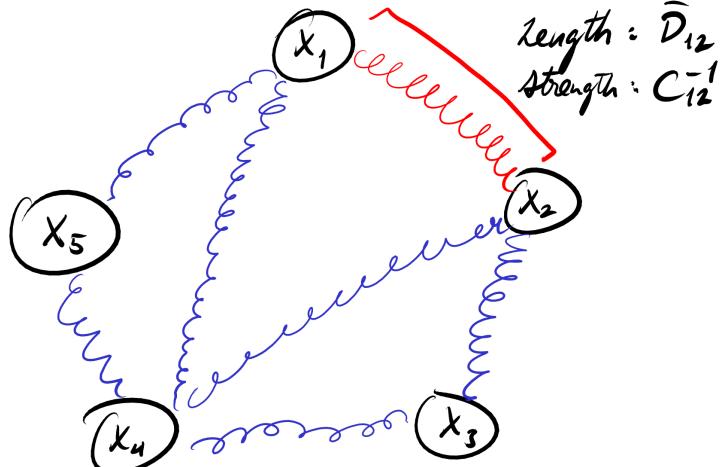


Pose Graph as Springs



$$\underset{\{X_1, \dots, X_n\}}{\text{minimize}} \quad \sum_{i,j} \left(\bar{D}_{ij} - (X_i - X_j) \right)^T C_{ij}^{-1} \left(\bar{D}_{ij} - (X_i - X_j) \right)$$

Pose Graph as Springs



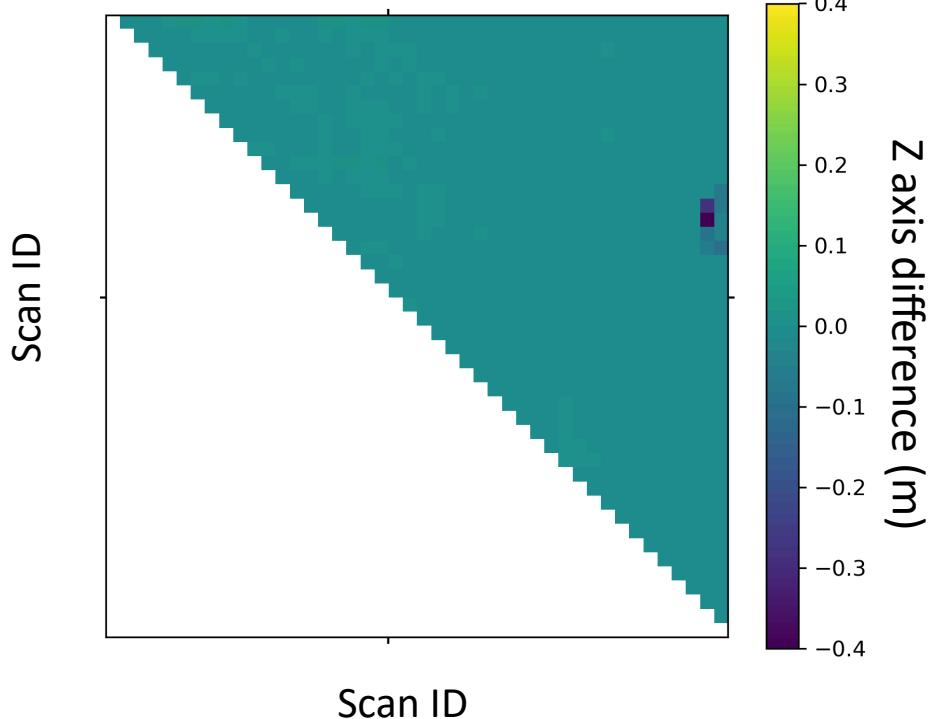
Ways to Set Spring Elasticities (covariances)

- Constant elasticity
- Model elasticity using pairwise features
 $(i, j) \rightarrow C_{ij}^{-1}$

$$\underset{\{X_1, \dots, X_n\}}{\text{minimize}} \quad \sum_{i,j} \left(\bar{D}_{ij} - (X_i - X_j) \right)^T C_{ij}^{-1} \left(\bar{D}_{ij} - (X_i - X_j) \right)$$

Rooting out Bad Pairwise Registrations

Comparing pairwise and global registrations



- After optimization, some springs (pairwise registrations) are stretched
- Lots of redundancy in the graph reveals poor pairwise registrations
- Weight those springs less in optimization step or remove them entirely

Lu-Milios Implementation

- Implemented registration optimizer using the closed-form least squares pose graph solution described in Lu and Milios' 1997 paper

$$\mathbf{x} = (\mathbf{H}^T \mathbf{C}^{-1} \mathbf{H})^{-1} \mathbf{H}^T \mathbf{C}^{-1} \bar{\mathbf{d}}, \quad \mathbf{C}_{\mathbf{x}} = (\mathbf{H}^T \mathbf{C}^{-1} \mathbf{H})^{-1}$$

Lu-Milios Implementation

- Implemented registration optimizer using the closed-form least squares pose graph solution described in Lu and Milios' 1997 paper

$$\mathbf{x} = (\mathbf{H}^T \mathbf{C}^{-1} \mathbf{H})^{-1} \mathbf{H}^T \mathbf{C}^{-1} \bar{\mathbf{d}}, \quad \mathbf{C}_{\mathbf{x}} = (\mathbf{H}^T \mathbf{C}^{-1} \mathbf{H})^{-1}$$

- Parameters:
 - ▶ Expected translational error of 1 meter, 2 mrad
(constant covariance of $\text{diag}(1, 1, 1, 0.002^2, 0.002^2, 0.002^2)$)
 - ▶ Prune edges (pairwise registrations) whose optimized Mahalanobis distance has z-score higher than 1.5 among all edges

Lu-Milios Implementation

- Implemented registration optimizer using the closed-form least squares pose graph solution described in Lu and Milios' 1997 paper

$$\mathbf{x} = (\mathbf{H}^T \mathbf{C}^{-1} \mathbf{H})^{-1} \mathbf{H}^T \mathbf{C}^{-1} \bar{\mathbf{d}}, \quad \mathbf{C}_{\mathbf{x}} = (\mathbf{H}^T \mathbf{C}^{-1} \mathbf{H})^{-1}$$

- Parameters:
 - Expected translational error of 1 meter, 2 mrad
(constant covariance of $\text{diag}(1, 1, 1, 0.002^2, 0.002^2, 0.002^2)$)
 - Prune edges (pairwise registrations) whose optimized Mahalanobis distance has z-score higher than 1.5 among all edges
- Note: Improper to perform linear least squares on Euler angles, but works alright since angles are very small (no more than a few mrad)

Comparing Old Optimizer with Lu-Milios

Optimization Time (Barrett Park)

Old ncc_nxn Optimizer

2m 54s

Lu-Milios Optimizer

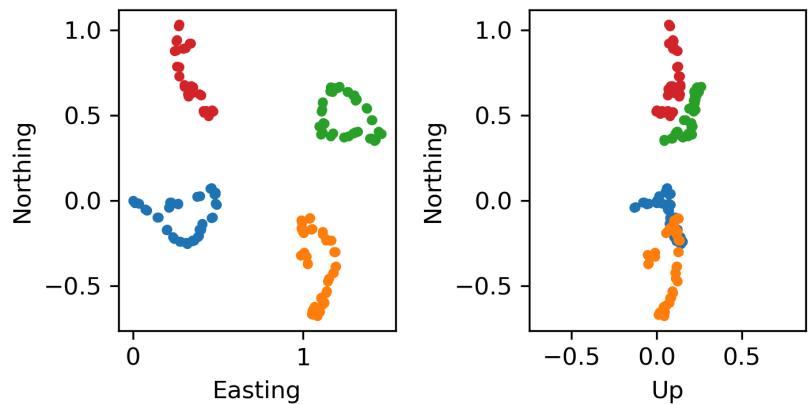
0m 3s

(58 times faster)

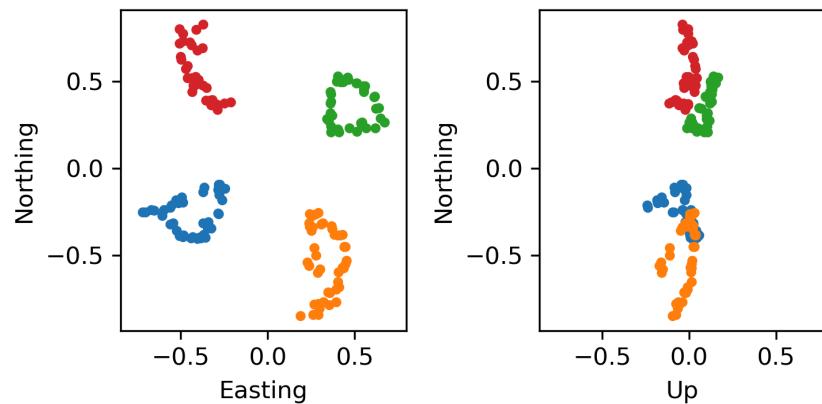
Comparing Old Optimizer with Lu-Milius

Optimized Translation Plots (Barrett Park)

Old ncc_nxn Optimizer



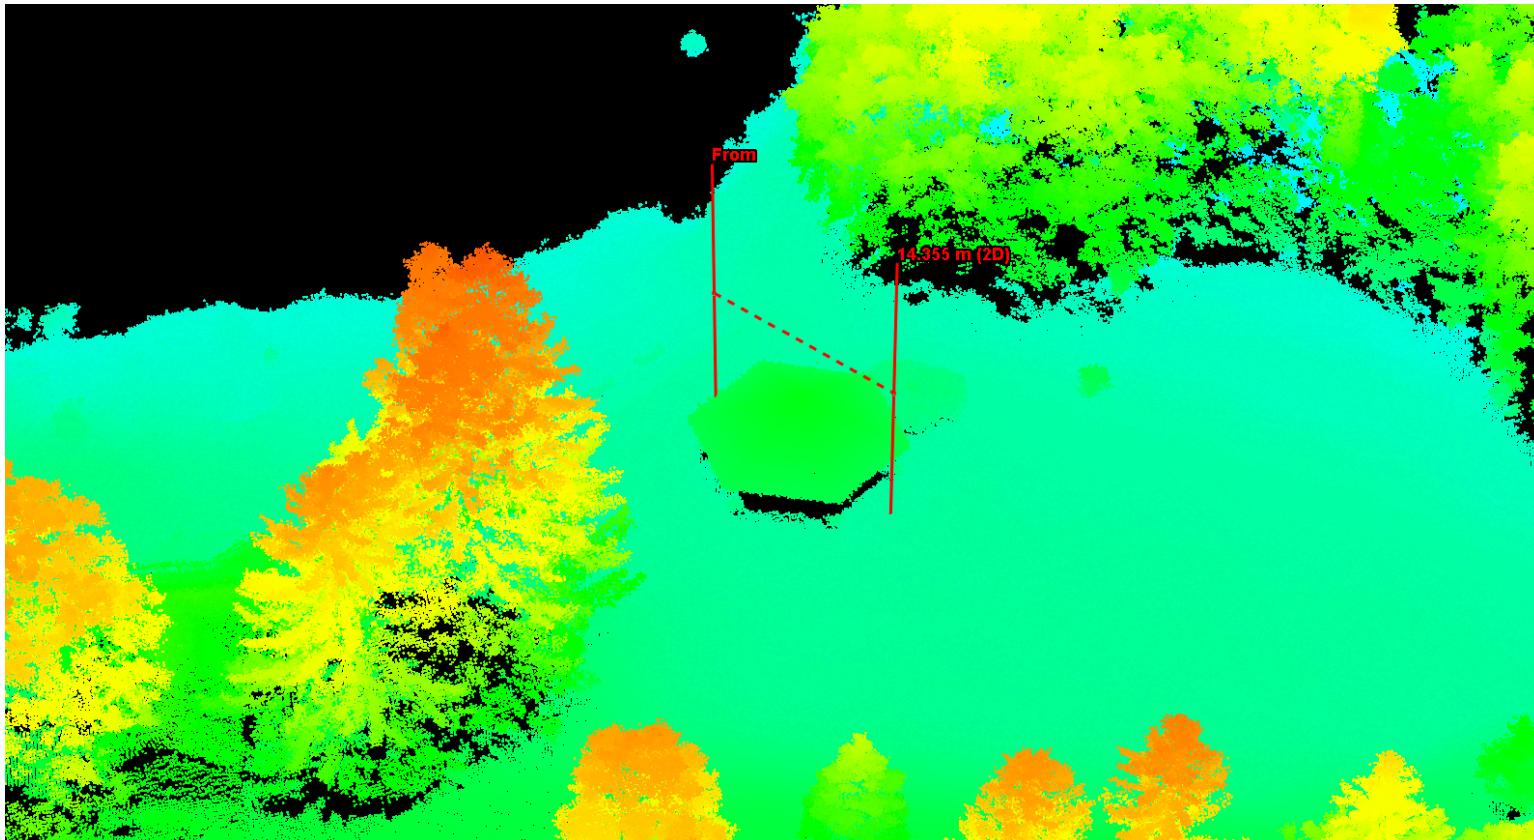
Lu-Milius Optimizer



Comparing Old Optimizer with Lu-Milios

Profile Analysis (Barrett Park, gazebo)

(only scans ending in scan00010)



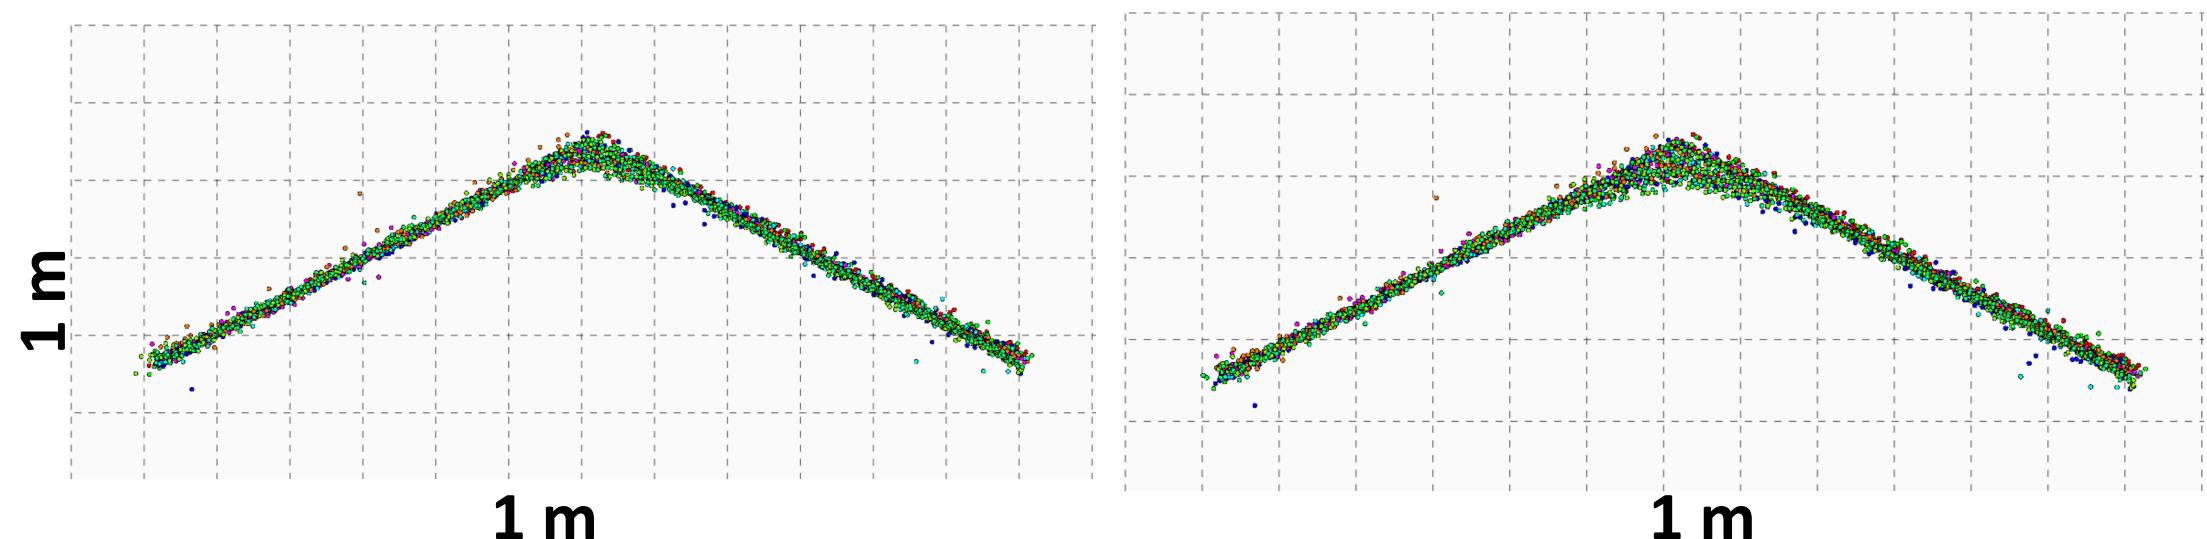
Comparing Old Optimizer with Lu-Milios

Profile Analysis (Barrett Park, gazebo)

(only scans ending in scan00010)

Old ncc_nxn Optimizer

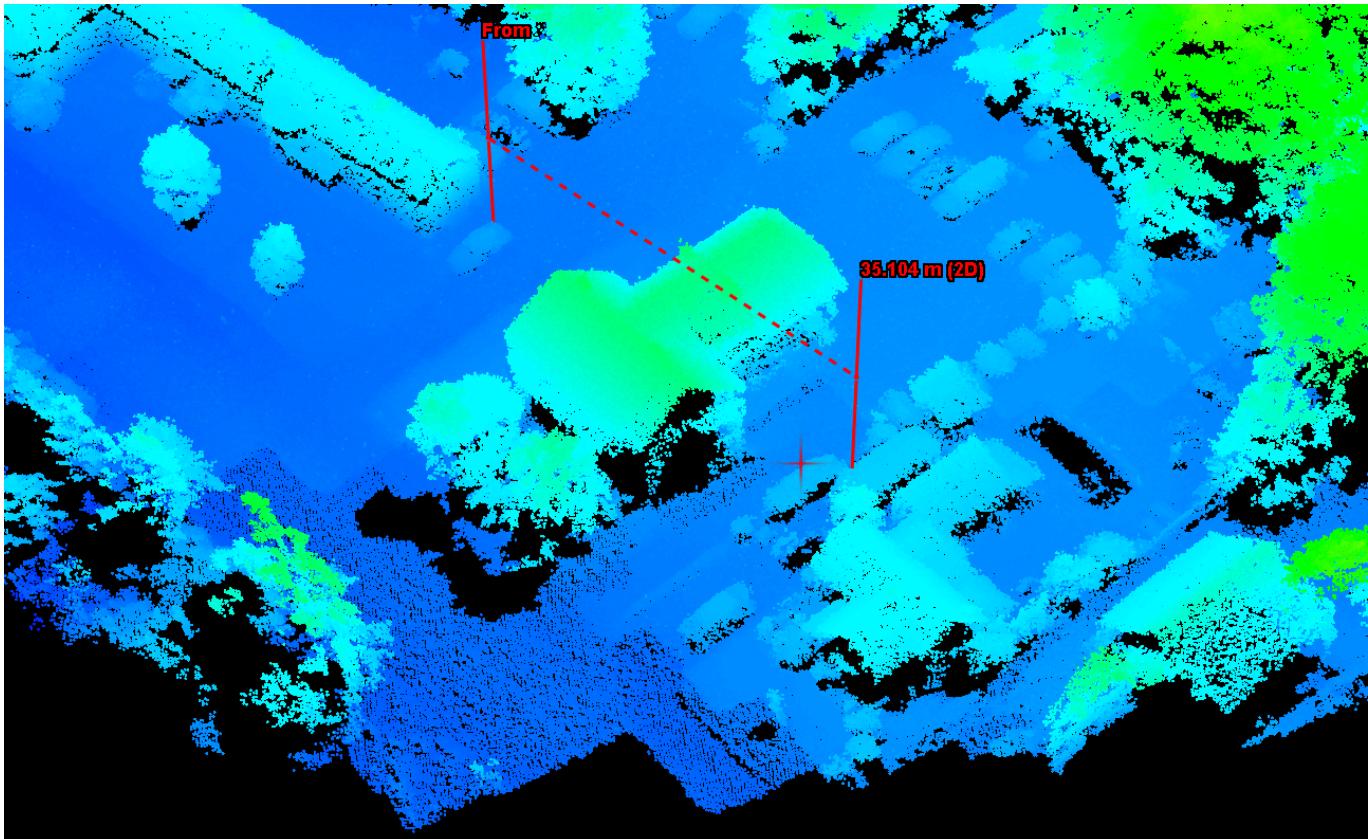
Lu-Milios Optimizer



Comparing Old Optimizer with Lu-Milios

Profile Analysis (Barrett Park, building on northeast)

(only scans ending in scan00010)

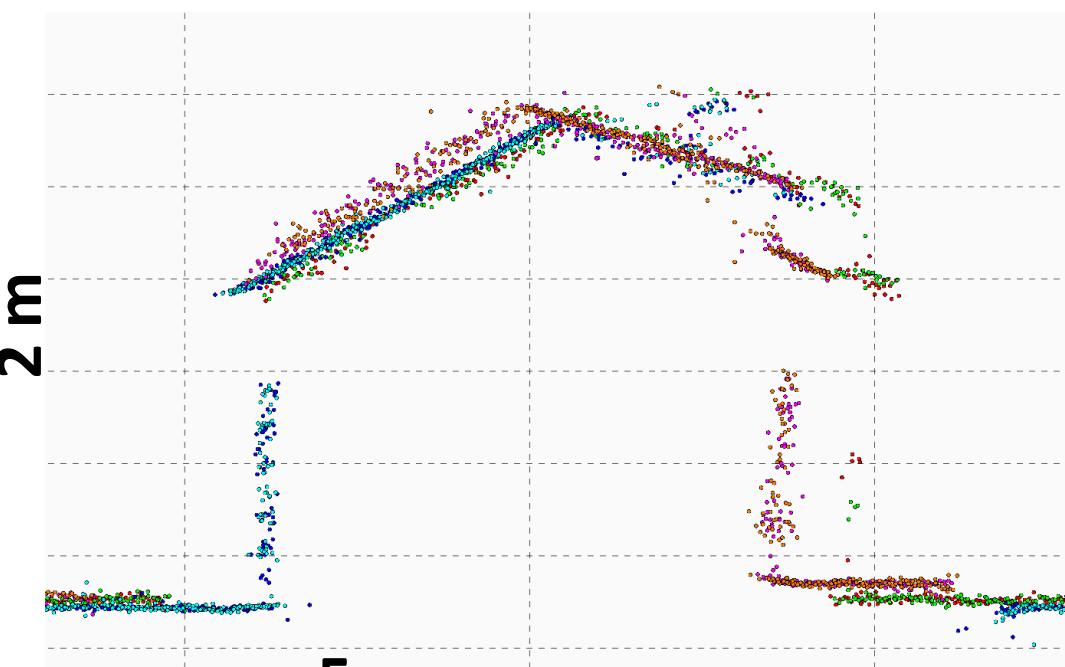


Comparing Old Optimizer with Lu-Milios

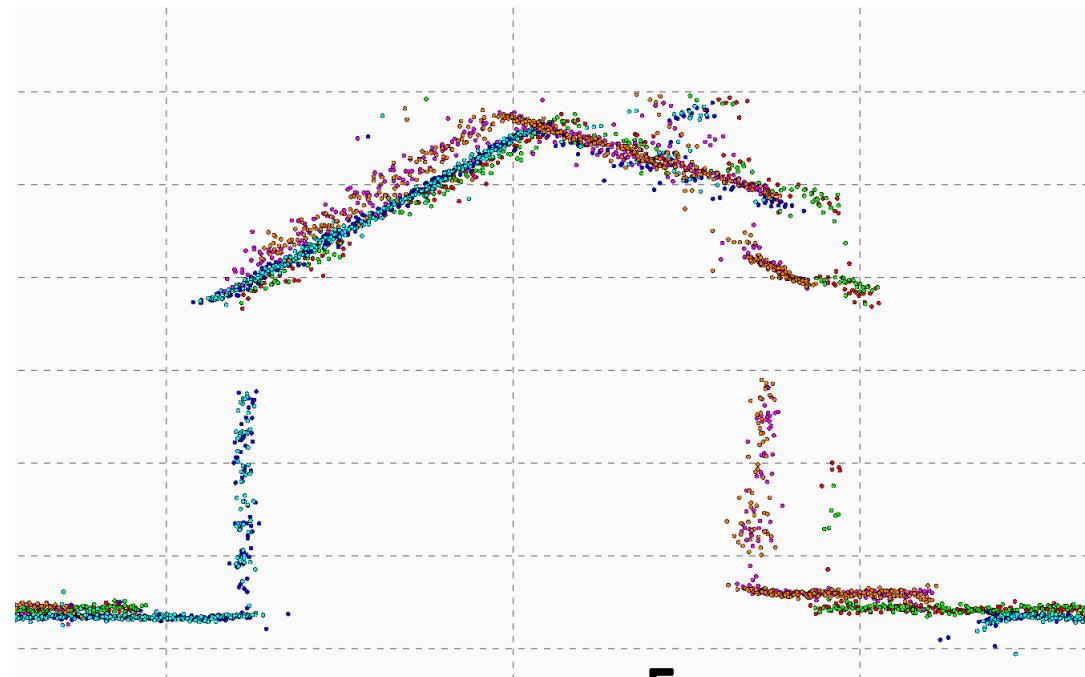
Profile Analysis (Barrett Park, building on northeast)

(only scans ending in scan00010)

Old ncc_nxn Optimizer



Lu-Milios Optimizer



(notice identical ghosting at upper left in both images)

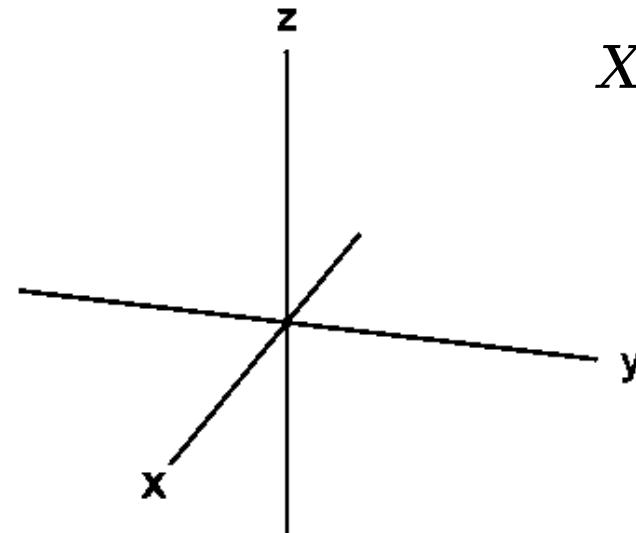
5 m

Ongoing and Future Work

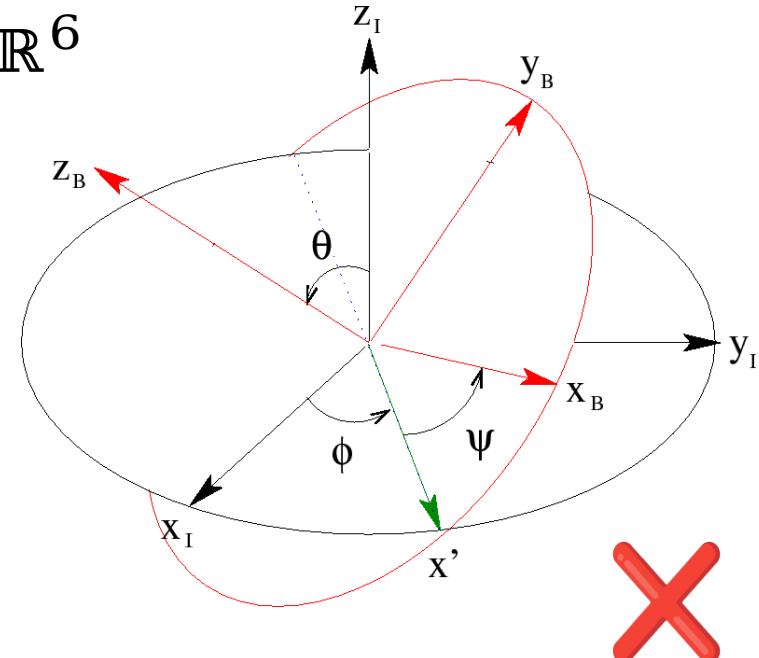
$$X_i = \begin{pmatrix} x_i \\ y_i \\ z_i \\ \theta_i \\ \phi_i \\ \psi_i \end{pmatrix} \in \mathbb{R}^6$$

$$\underset{\{X_1, \dots, X_n\}}{\text{minimize}} \quad \sum_{i,j} \left(\bar{D}_{ij} - (X_i - X_j) \right)^T C_{ij}^{-1} \left(\bar{D}_{ij} - (X_i - X_j) \right)$$

Ongoing and Future Work



$$X_i = \begin{pmatrix} x_i \\ y_i \\ z_i \\ \theta_i \\ \phi_i \\ \psi_i \end{pmatrix} \in \mathbb{R}^6$$



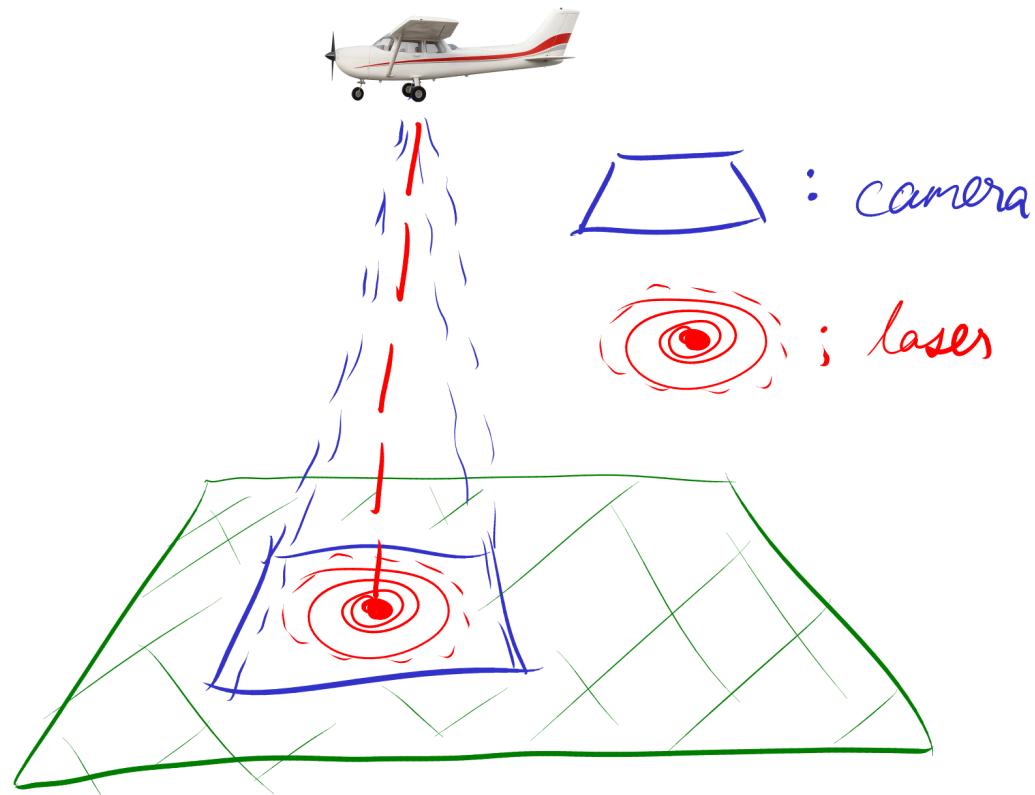
$$\underset{\{X_1, \dots, X_n\}}{\text{minimize}} \quad \sum_{i,j} \left(\bar{D}_{ij} - (X_i - X_j) \right)^T C_{ij}^{-1} \left(\bar{D}_{ij} - (X_i - X_j) \right)$$

Ongoing and Future Work

- **Use non-linear least squares optimization to better handle orientation**
Existing implementation spun its wheels for 3 hours on Barrett Park and failed...
- **Be smarter in choosing pairwise registration covariance (uncertainty)**
 - Automatically determine covariance using properties of the sensor or the data
 - Use pairwise registration features to determine covariance

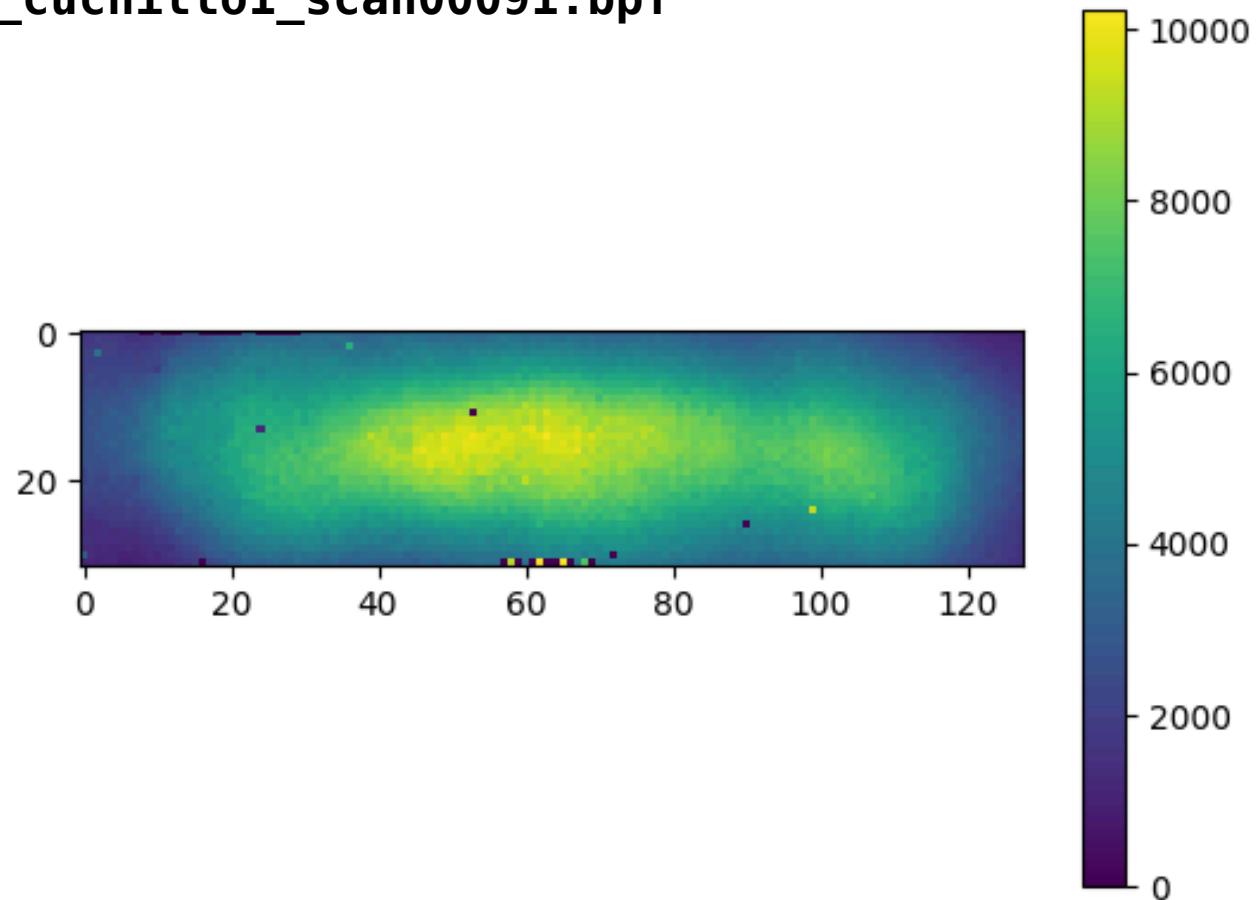
Modeling the Laser Illumination Spot

Laser Illumination Spot



Laser Illumination Spot

Total photon detections per pixel during
20230627_095732_cuchillo1_scan00091.bpf

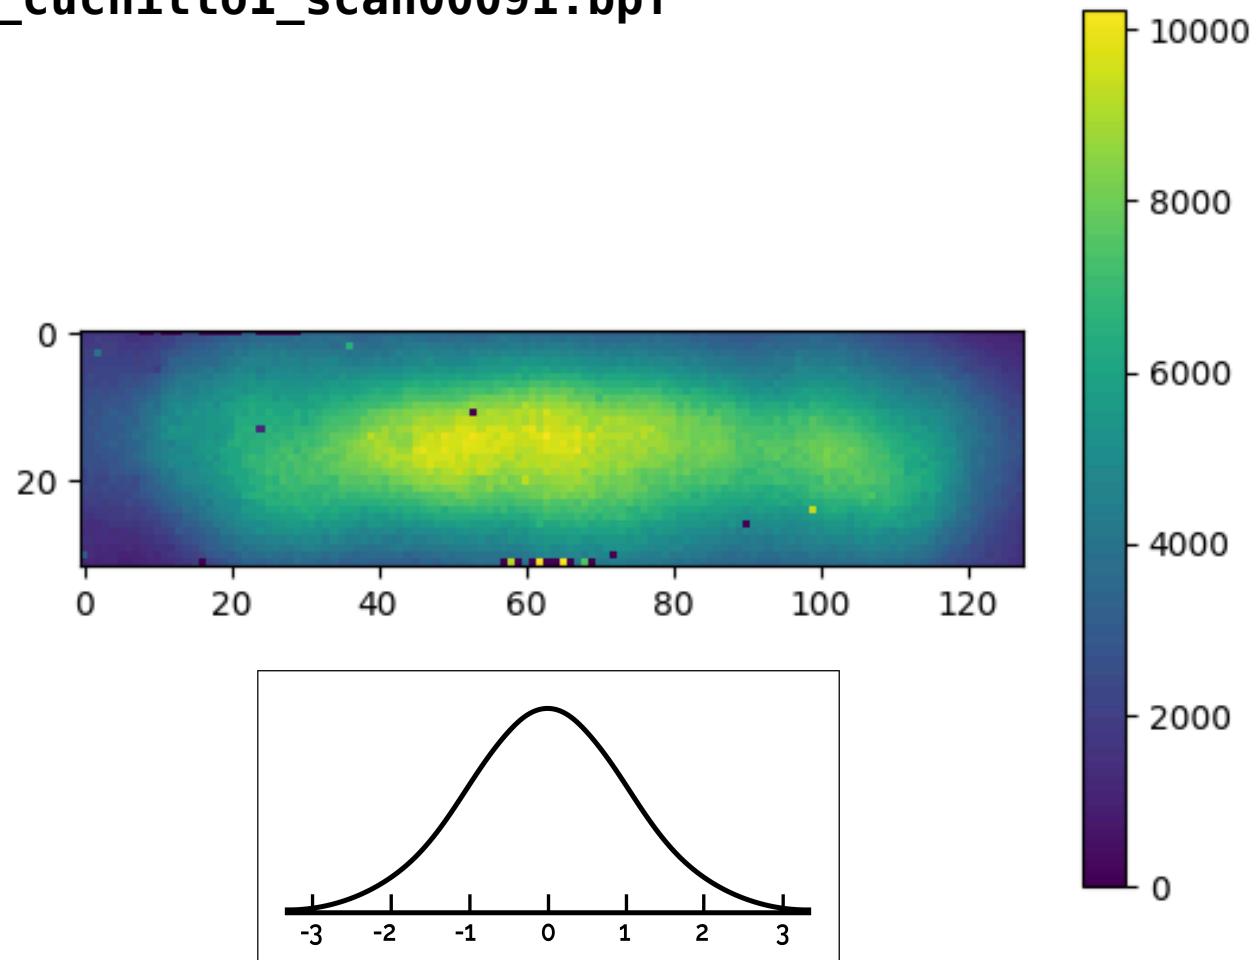


Why Model the Illumination Spot

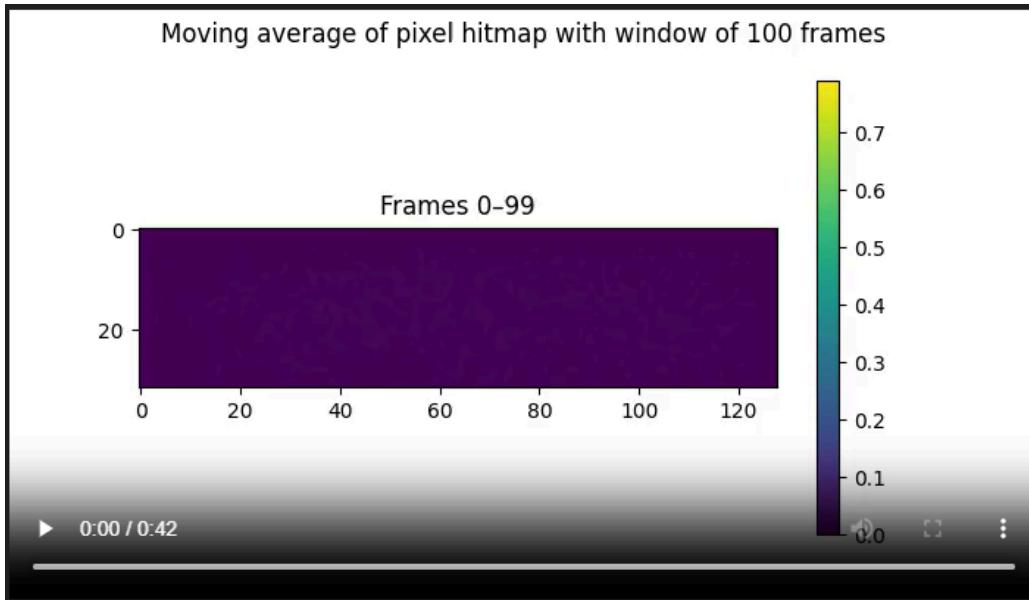
- More accurate pointwise reflectivity estimates
- Validate or refine alignment in-flight
- Track defective pixels

Spot is Approximately Gaussian

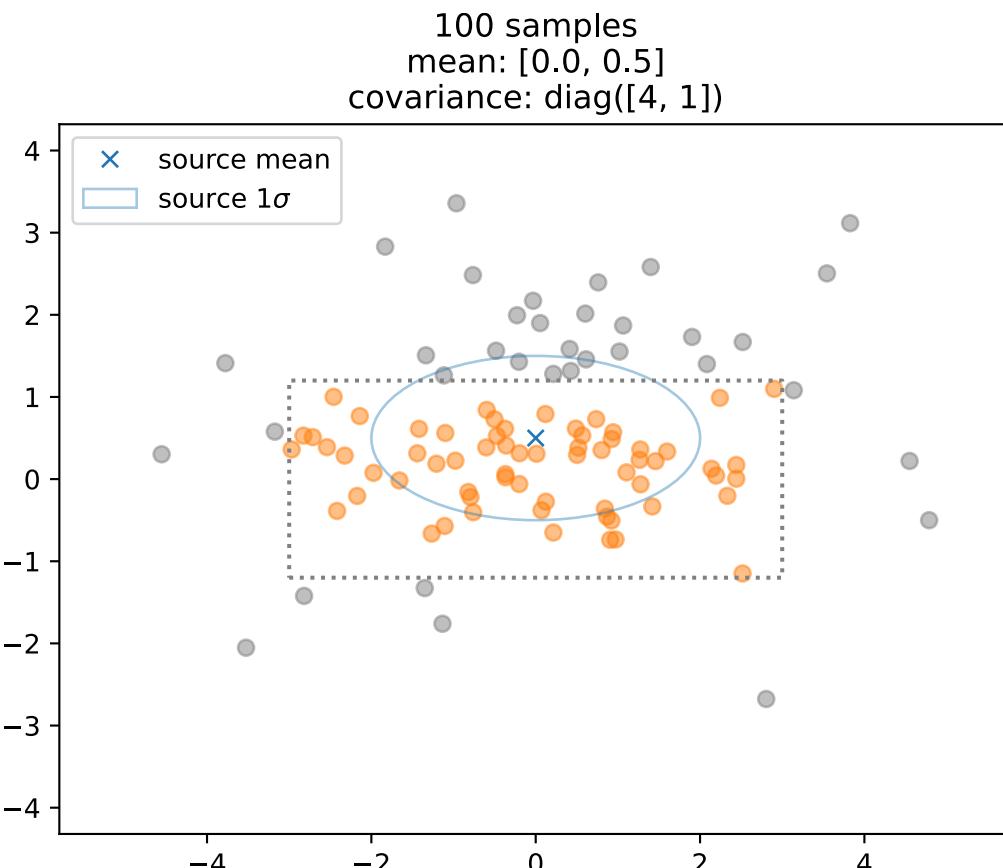
Total photon detections per pixel during
20230627_095732_cuchillo1_scan00091.bpf



Cannot Assume the Spot Stays Still

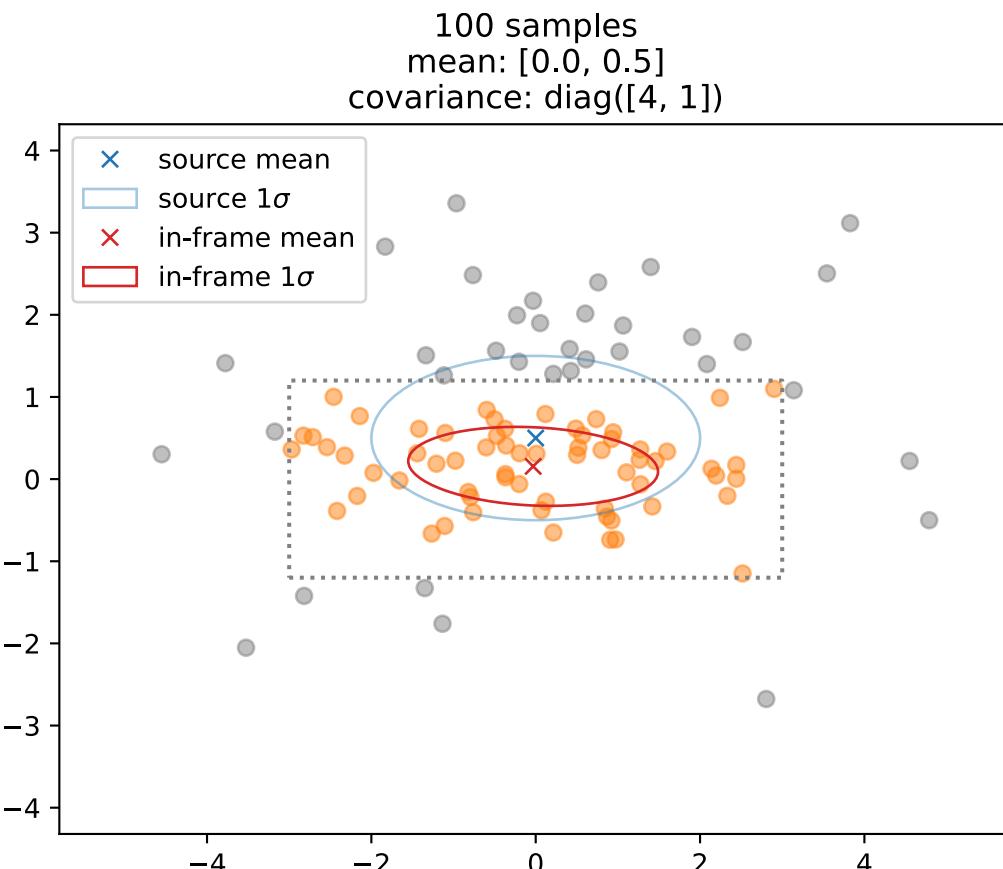


Cannot Use Sample Mean and Covariance



- 100 random Gaussian samples representing photons, not all of which are in-frame
- Ring represents 1 standard deviation
- We only see the samples within the frame

Cannot Use Sample Mean and Covariance



- 100 random Gaussian samples representing photons, not all of which are in-frame
- Ring represents 1 standard deviation
- We only see the samples within the frame
- Fitting Gaussian with sample mean and covariance doesn't work

Attempts

For each 100-frame pixel bitmap average, fit Gaussian parameters with:

- **Gaussian Mixture Models**

Dealt with problem just described—sample is occluded, means and covariances are off

Attempts

For each 100-frame pixel bitmap average, fit Gaussian parameters with:

- **Gaussian Mixture Models**

Dealt with problem just described—sample is occluded, means and covariances are off

- **Bayesian inference**

Necessary PyMC tools (cumulative distribution function of multivariate normal) [not fully implemented](#)

Attempts

For each 100-frame pixel bitmap average, fit Gaussian parameters with:

- **Gaussian Mixture Models**

Dealt with problem just described—sample is occluded, means and covariances are off

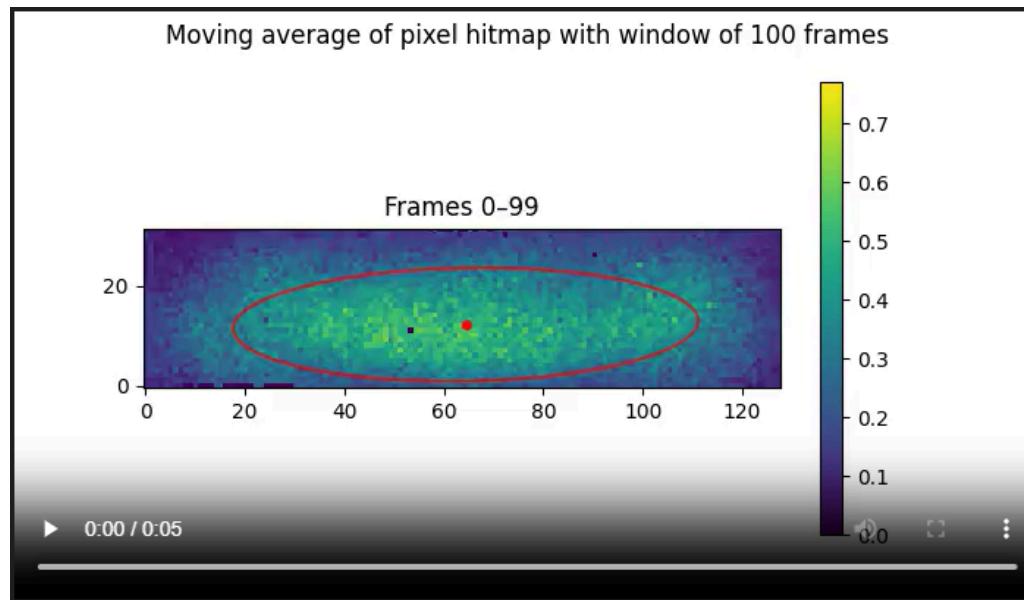
- **Bayesian inference**

Necessary PyMC tools (cumulative distribution function of multivariate normal) [not fully implemented](#)

- **pygmmis—Gaussian Mixture Models for occluded data**

- Written by Princeton astronomers to model luminosity of galaxies from truncated camera data (like ours)
- Gaussian Mixture Models (GMMs) that can handle occluded data
- Under the hood, uses Expectation Maximization (standard for GMMs), but also generates mock samples to handle occluded regions

Successful Single-Gaussian Fit



Future Work

- Continue investigating faster alternatives, using pygmmis as ground truth
 - Blur pixel heatmap to quickly track spot center without having to recalculate covariance (if single Gaussian)
 - Fit to some subset of the data—random subset works surprisingly well
- Test on more scans
- Implement reflectivity estimate normalized by spot illumination model

References

Processing Performance Report

- [processing-performance \(master\) on Bitbucket](#)
- [acadia \(master\) on Bitbucket](#) Apptainers in apptainer_creation and slurm bricks in processing_workflow
- [python_3deo/fileIO/readGSOF.py](#) Used to collect certain scanning information

Pose Graphs for Registration

- [Presentation I gave on pose graph registration](#)
- [Ideas related to the above presentation](#)
- [Lu and Miliros paper](#) “Globally Consistent Range Scan Alignment for Environment Mapping”, April 1997. Introduces these ideas in the context of robotics
- [Old optimizer](#) Replaced by Lu-Miliros in zreg_ncc commit 2b6d0d4 on July 18th
- [Lu-Miliros implementation](#) zreg_ncc, python/lu_miliros.py

Modeling the Laser Illumination Spot

- [Write-up](#) Describes background and my work thus far on this
- [PyMC forum discussion](#) I asked about using Bayesian inference—necessary tools not implemented